

In [101]:

```
#Creator: Hetul Varaiya  
#University: Arizona State University  
#Date: 04/13/2018 - 04/14/2018  
#I am an industrial Engineering Graduate student with a deep interest in data analytics  
and machine Learning  
  
#THE APPROACH USED INCLUDES:  
#SVM - LINEAR SVC - 76%  
#LOGISTIC REGRESSION - 71%  
#NEAREST NEIGHBOR - 60%  
#XGBOOST  
#RANDOM FOREST
```

In [2]:

```
#IMPORTING LIBRARY  
  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
%matplotlib inline
```

In [3]:

```
# PREPROCESSING  
#FILLING NA VALUES  
#TRANSFORMING DATA  
  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import xgboost as xgb  
from sklearn import preprocessing  
  
dftr = pd.read_csv('K:/Mckinsey hackathon/without pca/train.csv')  
stro1 = dftr['stroke']  
del dftr['stroke']  
dftst = pd.read_csv('K:/Mckinsey hackathon/without pca/test.csv')  
idkey = dftst['id']  
  
del dftr['id']  
del dftst['id']  
dftr.isnull().any()  
dftst.isnull().any()  
dftr[['smoking_status']] = dftr[['smoking_status']].fillna(value='hetul')  
dftst[['smoking_status']] = dftst[['smoking_status']].fillna(value='hetul')  
  
from sklearn import preprocessing
```

In [4]:

```
#CREATED DUMMIES TO MAKE IT SINGLE-CLASS[0,1]
p = pd.get_dummies(data= dftr, columns = ['smoking_status','ever_married','Residence_type','work_type','gender'])
del p['smoking_status_hetul']

q = pd.get_dummies(data= dftst, columns = ['smoking_status','ever_married','Residence_type','work_type','gender'])
del q['smoking_status_hetul']
```

In [5]:


```
# TAKING THE MEAN BECAUSE VERY FEW VALUES OF BMI ARE NAN
p[['bmi']] = p[['bmi']].fillna(p[['bmi']].mean())
q[['bmi']] = q[['bmi']].fillna(q[['bmi']].mean())
```

In [6]:

```
#FINDING CORRELATION TO GET AN OVERVIEW OF THE RELATION BETWEEN THE COLUMNS
corr= dftr.corr()
print(np.matrix(corr))
mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True
f, ax = plt.subplots(figsize=(13, 13))

cmap = sns.diverging_palette(225, 12, as_cmap=True)
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=1, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})
plt.show(sns.heatmap(corr, mask=mask, cmap=cmap, vmax=1, center=0,
                    square=True, linewidths=.5, cbar_kws={"shrink": .5}))
```

```
[[1.          0.27216879 0.25018839 0.23762684 0.35889719]
 [0.27216879 1.          0.11977703 0.16021129 0.16122524]
 [0.25018839 0.11977703 1.          0.14693807 0.05767724]
 [0.23762684 0.16021129 0.14693807 1.          0.19129515]
 [0.35889719 0.16122524 0.05767724 0.19129515 1.          ]]
```



In [7]:

```
#USED STANDARDSCALER TO TRANSFORM THE DATA  
from sklearn.preprocessing import StandardScaler  
  
scaler = StandardScaler()  
scaler.fit(p)  
  
scaler.mean_  
p = scaler.transform(p)  
  
scaler.fit(q)  
  
scaler.mean_  
q = scaler.transform(q)
```

In [102]:

```
# PCA  
  
from sklearn.decomposition import PCA  
pca = PCA(n_components=2)  
p = pca.fit_transform(p)  
p = pd.DataFrame(data = p)  
q= pca.fit_transform(q)  
q = pd.DataFrame(data = q)
```

In [10]:

```
p2 = p.copy()
```

In [11]:

```
X = p2  
y = stro1
```

In [12]:

```
#splitting the data for validation  
  
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=  
42)
```

In [13]:

```
#FUNCTION USED FOR ROC VALUES AND CURVE PLOT
def ROC(label,result):
    # Compute ROC curve and area the curve
    Y = np.array(label)
    fpr, tpr, thresholds = roc_curve(Y, result)
    roc_auc = auc(fpr, tpr)
    print("Area under the ROC curve : %f" % roc_auc)

    # Plot ROC curve
    #pl.clf()
    plt.plot(fpr, tpr, label='ROC curve (area = %0.2f)' % roc_auc)
    plt.plot([0, 1], [0, 1], 'k--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.0])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver operating characteristic')
    plt.legend(loc="lower right")
    plt.show()
```

In [62]:

```
#svm using L2 regularization using the training and testing data to do the validation

from sklearn.utils import shuffle
from sklearn.metrics import roc_curve, auc
from sklearn import svm
import numpy as np
import itertools
import matplotlib.pyplot as plt
from sklearn.utils import class_weight
from sklearn.model_selection import GridSearchCV
from sklearn.datasets import make_moons, make_circles, make_classification

linearly_separable = (X_train, y_train)
Cs = [0.001, 0.01,0.1,1]
parameters = {'C':Cs}

# svc2tr = svm.SVC(class_weight='balanced',probability = True)
# clfl2tr = GridSearchCV(svc2tr, parameters,return_train_score=True)
# clfl2tr.fit(X_train, y_train)
# # y_predsvml2tr = clfl2tr.predict(X_test)
# #print(y_predsvml2)

# # y_score = clfl2.predict_proba(X_test)[: ,1]

# # ROC(y_test,y_score)
# probl2tr = clfl2tr.predict_proba(X_test)
# y_scorel2tr = clfl2tr.predict_proba(X_test)[: ,1]
```

In [98]:

```
from sklearn.svm import LinearSVC
from sklearn.calibration import CalibratedClassifierCV
svml2tr = LinearSVC(class_weight='balanced')
clf12tr = CalibratedClassifierCV(svml2tr)
clf12tr.fit(X_train, y_train)
probl2tr = clf12tr.predict_proba(X_test)
y_scorel2tr = clf12tr.predict_proba(X_test)[:,-1]
```

In [16]:

```
# from sklearn.metrics import roc_curve
# roc = roc_curve(y_predsvml2, y_test)
# print(roc)
```

In [17]:


```
# from sklearn.cross_validation import cross_val_score
# cross_val_score(svc2tr, X_train, y_train, scoring = 'roc_auc')
```

In [99]:

```
#ROC and AUC measure using the ROC function

ROC(y_test,y_scorel2tr)
```

Area under the ROC curve : 0.790681



In [100]:

```
#svm using L2 regularization for the actual data
# SVC seems to give the best output out of all the models tried after tuning the parameters
# from sklearn.model_selection import train_test_split
# X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)
from sklearn import svm
from sklearn.metrics import confusion_matrix
import numpy as np
import itertools
import matplotlib.pyplot as plt
from sklearn.utils import class_weight
# from sklearn import preprocessing
# from sklearn.metrics import accuracy_score
from sklearn.model_selection import GridSearchCV
from sklearn.datasets import make_moons, make_circles, make_classification

# n_samples=300
# linearly_separable = (X, y)
# Cs = [0.001, 0.01,0.1,1]
# parameters = {'C':Cs}

from sklearn.svm import LinearSVC
from sklearn.calibration import CalibratedClassifierCV
svml2 = LinearSVC(class_weight='balanced')
clfl2 = CalibratedClassifierCV(svml2)
clfl2.fit(X, y)
y_scorel2tr = clf.predict_proba(q)[:,-1]
# svc2 = svm.LinearSVC(penalty = 'l2',class_weight='balanced')
# clfl2 = GridSearchCV(svc2, parameters,return_train_score=True)
# clfl2.fit(X, y)
# probl2 = clfl2.predict_proba(q)
# y_scorel2 = clfl2.predict_proba(q)[:,-1]
# print(y_scorel2)

#np.savetxt('K:/svml2.txt',y_scorel2, delimiter="\t")

# # print('best parameters for l2:',clfl2.best_params_,',testing accuracy for l2:',clfl2.best_estimator_.score(y_pred,y_test))
# # from sklearn.metrics import classification_report
# # print(classification_report(y_test,y_pred))
# # cm = confusion_matrix(y_test, y_pred)
# # p = plot_confusion_matrix(cm,[0,1],normalize=True)
```

In [97]:

```
np.savetxt('K:/svml2.txt',y_scorel2tr, delimiter="\t")
```


In [40]:

```
#logistic regression using the training and testing data to do the validation

import sklearn.linear_model as skl_lm
import seaborn as sns
from sklearn.utils import shuffle
from sklearn.metrics import roc_curve, auc

clflrtr = skl_lm.LogisticRegression(class_weight='balanced', solver = 'lbfgs')
clflrtr.fit(X_train,y_train)
# prob = clflr.predict_proba(q)
y_scorelrtr = clflrtr.predict_proba(X_test)[: ,1]

# y_predlrtr = clflrtr.predict(X_test)
# print(y_predlrtr)


#np.savetxt('K:/logistic regression.txt',np.c_[idkey,y_predlr], delimiter="\t")
# np.savetxt('myfile.txt', np.c_[x,y,z])
```

In [42]:

```
#ROC and AUC measure using the ROC function

ROC(y_test,y_scorelrtr)
```

Area under the ROC curve : 0.791473



In [43]:

```
#Logistic regression for the actual data

import sklearn.linear_model as skl_lm
from sklearn.metrics import average_precision_score
import seaborn as sns
from sklearn.utils import shuffle
from sklearn.metrics import roc_curve, auc

clflr = skl_lm.LogisticRegression(class_weight='balanced', solver = 'lbfgs')
clflr.fit(X,y)
# prob = clflr.predict_proba(q)
# y_predlr = clflr.predict(q)
# print(y_predlr)
problr = clflr.predict_proba(q)
y_scorelr = clflr.predict_proba(q)[:,:1]

#np.savetxt('K:/logistic regression.txt',y_scorelr, delimiter="\t")
# np.savetxt('myfile.txt', np.c_[x,y,z])
```

In [84]:

```
#Random forest using the training and testing data to do the validation

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
# from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import classification_report
# from sklearn.metrics import matthews_corrcoef
from sklearn.metrics import roc_curve, auc

depthsi = [1,2,4,8,16,32,64]
ntrees = [1,2,4,8,16,32,64]
max_features = [0.1,0.2,0.3,0.5,0.9]

parameters = {'max_depth':depthsi, 'n_estimators':ntrees, 'max_features':max_features}

rfctr = RandomForestClassifier(class_weight='balanced',max_depth=depthsi,n_estimators =
    ntrees, max_features = max_features )
clfrftr = GridSearchCV(rfctr, parameters,return_train_score=True)
clfrftr.fit(X_train, y_train)
print('best parameters:',clfrftr.best_params_)


best parameters: {'max_depth': 32, 'max_features': 0.5, 'n_estimators': 6
4}
```

In [85]:

```
y_scorerftr = clfrftr.predict_proba(X_test)[:,-1]
#print(y_predrf)

#ROC and AUC measure using the ROC function
ROC(y_test,y_scorerftr)
```

Area under the ROC curve : 0.631577



In [86]:

```
#Random forest for the actual data

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
# from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import classification_report
# from sklearn.metrics import matthews_corrcoef
from sklearn.metrics import roc_curve, auc

depthsi = [1,2,4,8,16,32,64]
ntrees = [1,2,4,8,16,32,64]
max_features = [0.1,0.2,0.3,0.5,0.9]

parameters = {'max_depth':depthsi, 'n_estimators':ntrees, 'max_features':max_features}

rfc = RandomForestClassifier(class_weight='balanced',max_depth=depthsi,n_estimators = n
trees, max_features = max_features )
clfrf = GridSearchCV(rfc, parameters,return_train_score=True)
clfrf.fit(X, y)
print('best parameters:',clfrf.best_params_)
```

```

-----
----
KeyboardInterrupt                                Traceback (most recent call 1
ast)
<ipython-input-86-9556c8897df2> in <module>()
    17 rfc = RandomForestClassifier(class_weight='balanced',max_depth=
depthsi,n_estimators = ntrees, max_features = max_features )
    18 clfrf = GridSearchCV(rfc, parameters,return_train_score=True)
--> 19 clfrf.fit(X, y)
    20 print('best parameters:',clfrf.best_params_)
    21

~\Anaconda3\lib\site-packages\sklearn\model_selection\_search.py in fit
(self, X, y, groups, **fit_params)
    637                                     error_score=self.error_score)
    638         for parameters, (train, test) in product(candidate_pa
rams,
--> 639                                     cv.split(X,
y, groups)))
    640
    641         # if one choose to see train score, "out" will contain
train score info

~\Anaconda3\lib\site-packages\sklearn\externals\joblib\parallel.py in _
_call__(self, iterable)
    777         # was dispatched. In particular this covers the edg
e
    778         # case of Parallel used with an exhausted iterator.
--> 779         while self.dispatch_one_batch(iterator):
    780             self._iterating = True
    781         else:

~\Anaconda3\lib\site-packages\sklearn\externals\joblib\parallel.py in d
ispatch_one_batch(self, iterator)
    623             return False
    624         else:
--> 625             self._dispatch(tasks)
    626             return True
    627

~\Anaconda3\lib\site-packages\sklearn\externals\joblib\parallel.py in _
dispatch(self, batch)
    586         dispatch_timestamp = time.time()
    587         cb = BatchCompletionCallBack(dispatch_timestamp, len(ba
tch), self)
--> 588         job = self._backend.apply_async(batch, callback=cb)
    589         self._jobs.append(job)
    590

~\Anaconda3\lib\site-packages\sklearn\externals\joblib\_parallel_backen
ds.py in apply_async(self, func, callback)
    109     def apply_async(self, func, callback=None):
    110         """Schedule a func to be run"""
--> 111         result = ImmediateResult(func)
    112         if callback:
    113             callback(result)

~\Anaconda3\lib\site-packages\sklearn\externals\joblib\_parallel_backen
ds.py in __init__(self, batch)
    330         # Don't delay the application, to avoid keeping the inp
ut

```

```

331         # arguments in memory
--> 332         self.results = batch()
333
334     def get(self):

~\Anaconda3\lib\site-packages\sklearn\externals\joblib\parallel.py in _
__call__(self)
129
130     def __call__(self):
--> 131         return [func(*args, **kwargs) for func, args, kwargs in
self.items]
132
133     def __len__(self):

~\Anaconda3\lib\site-packages\sklearn\externals\joblib\parallel.py in <
listcomp>(.0)
129
130     def __call__(self):
--> 131         return [func(*args, **kwargs) for func, args, kwargs in
self.items]
132
133     def __len__(self):

~\Anaconda3\lib\site-packages\sklearn\model_selection\_validation.py in
_fit_and_score(estimator, X, y, scorer, train, test, verbose, paramete
rs, fit_params, return_train_score, return_parameters, return_n_test_sa
mples, return_times, error_score)
490         if return_train_score:
491             train_scores = _score(estimator, X_train, y_train,
scorer,
--> 492                                     is_multimetric)
493
494         if verbose > 2:

~\Anaconda3\lib\site-packages\sklearn\model_selection\_validation.py in
_score(estimator, X_test, y_test, scorer, is_multimetric)
521         """
522         if is_multimetric:
--> 523             return _multimetric_score(estimator, X_test, y_test, sc
orer)
524         else:
525             if y_test is None:

~\Anaconda3\lib\site-packages\sklearn\model_selection\_validation.py in
_multimetric_score(estimator, X_test, y_test, scorers)
551         score = scorer(estimator, X_test)
552         else:
--> 553             score = scorer(estimator, X_test, y_test)
554
555         if hasattr(score, 'item'):

~\Anaconda3\lib\site-packages\sklearn\metrics\scorer.py in _passthrough
_scorer(estimator, *args, **kwargs)
242 def _passthrough_scorer(estimator, *args, **kwargs):
243     """Function that wraps estimator.score"""
--> 244     return estimator.score(*args, **kwargs)
245
246

~\Anaconda3\lib\site-packages\sklearn\base.py in score(self, X, y, samp
le_weight)

```

```

347         """
348         from .metrics import accuracy_score
--> 349         return accuracy_score(y, self.predict(X), sample_weight
=sample_weight)
350
351

```

```

~\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py in predict(self, X)

```

```

536         The predicted classes.
537         """
--> 538         proba = self.predict_proba(X)
539
540         if self.n_outputs_ == 1:

```

```

~\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py in predict_proba(self, X)

```

```

587         Parallel(n_jobs=n_jobs, verbose=self.verbose, backend
="threading")(
588             delayed(accumulate_prediction)(e.predict_proba, X,
all_proba, lock)
--> 589             for e in self.estimators_)
590
591         for proba in all_proba:

```

```

~\Anaconda3\lib\site-packages\sklearn\externals\joblib\parallel.py in _call__(self, iterable)

```

```

777         # was dispatched. In particular this covers the edge
e
778         # case of Parallel used with an exhausted iterator.
--> 779         while self.dispatch_one_batch(iterator):
780             self._iterating = True
781         else:

```

```

~\Anaconda3\lib\site-packages\sklearn\externals\joblib\parallel.py in dispatch_one_batch(self, iterator)

```

```

623         return False
624         else:
--> 625         self._dispatch(tasks)
626         return True
627

```

```

~\Anaconda3\lib\site-packages\sklearn\externals\joblib\parallel.py in _dispatch(self, batch)

```

```

586         dispatch_timestamp = time.time()
587         cb = BatchCompletionCallBack(dispatch_timestamp, len(batch), self)
--> 588         job = self._backend.apply_async(batch, callback=cb)
589         self._jobs.append(job)
590

```

```

~\Anaconda3\lib\site-packages\sklearn\externals\joblib\_parallel_backend.py in apply_async(self, func, callback)

```

```

109         def apply_async(self, func, callback=None):
110             """Schedule a func to be run"""
--> 111             result = ImmediateResult(func)
112             if callback:
113                 callback(result)

```

```

~\Anaconda3\lib\site-packages\sklearn\externals\joblib\_parallel_backend.py in __init__(self, batch)

```

```

330      # Don't delay the application, to avoid keeping the inp
ut
331      # arguments in memory
--> 332      self.results = batch()
333
334      def get(self):

~\Anaconda3\lib\site-packages\sklearn\externals\joblib\parallel.py in _
__call__(self)
129
130      def __call__(self):
--> 131      return [func(*args, **kwargs) for func, args, kwargs in
self.items]
132
133      def __len__(self):

~\Anaconda3\lib\site-packages\sklearn\externals\joblib\parallel.py in <
listcomp>(.0)
129
130      def __call__(self):
--> 131      return [func(*args, **kwargs) for func, args, kwargs in
self.items]
132
133      def __len__(self):

~\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py in accumulate_
prediction(predict, X, out, lock)
381
382 def accumulate_prediction(predict, X, out, lock):
--> 383     prediction = predict(X, check_input=False)
384     with lock:
385         if len(out) == 1:

~\Anaconda3\lib\site-packages\sklearn\tree\tree.py in predict_proba(sel
f, X, check_input)
820     check_is_fitted(self, 'tree_')
821     X = self._validate_X_predict(X, check_input)
--> 822     proba = self.tree_.predict(X)
823
824     if self.n_outputs_ == 1:

```

KeyboardInterrupt:

In [218]:

```

# y_predrf = (clfrf.predict_proba(X_test)>0.5)[: ,1]
y_scorerf = clfrf.predict_proba(q)[: ,1]
# print(y_predrf)

np.savetxt('K:/Mckinsey hackathon/random forest.txt',y_scorerf, delimiter="\t")

[0 0 0 ... 0 0 0]

```


In [28]:


```
#xgboost using the training and testing data to do the validation
import warnings
warnings.filterwarnings("ignore")
from xgboost import XGBClassifier
depthsi = [1,2,4,8,16]
ntrees = [1,2,4,8,16]
max_features = [0.1,0.2,0.3,0.5,0.9]

parameters = {'max_depth':depthsi, 'n_estimators':ntrees, 'max_features':max_features}
gbttr = XGBClassifier(max_depth=depthsi,n_estimators = ntrees, max_features = max_features,
class_weight='balanced' )
clfxbtr = GridSearchCV(gbttr, parameters,return_train_score=True)
clfxbtr.fit(X, y)
print('best parameters:',clfxbtr.best_params_)
# y_score = clf.predict_proba(q)[:,:1]

# y_pred = (clf.predict_proba(q)>0.5)[:,:1]
y_predxbtr = clfxbtr.predict(X_test)
ROC(y_test,y_predxbtr)

# np.savetxt('K:/Mckinsey hackathon/xgboost.txt',y_predxbtr, delimiter="\t")
```

best parameters: {'max_depth': 1, 'max_features': 0.1, 'n_estimators': 1}
Area under the ROC curve : 0.500000



In [29]:

```
#xgboost for the actual data

import warnings
warnings.filterwarnings("ignore")
from xgboost import XGBClassifier
depthsi = [1,2,4,8,16]
ntrees = [1,2,4,8,16]
max_features = [0.1,0.2,0.3,0.5,0.9]

parameters = {'max_depth':depthsi, 'n_estimators':ntrees, 'max_features':max_features}
gbt = XGBClassifier(max_depth=depthsi,n_estimators = ntrees, max_features = max_features, class_weight='balanced' )
clf_xgb = GridSearchCV(gbt, parameters, return_train_score=True)
clf_xgb.fit(X, y)
print('best parameters:', clf_xgb.best_params_)
# y_score = clf.predict_proba(q)[: ,1]

# y_pred = (clf.predict_proba(q)>0.5)[: ,1]
# ROC(y_test,y_score)
y_pred_xgb = clf_xgb.predict(q)

np.savetxt('K:/Mckinsey hackathon/xgboost.txt', y_pred_xgb, delimiter="\t")
```

best parameters: {'max_depth': 1, 'max_features': 0.1, 'n_estimators': 1}

In [206]:

```
# print(y_pred_xgb)
```


In [92]:

```
from sklearn import neighbors
n_neighbors = 15
clfnn = neighbors.KNeighborsClassifier(n_neighbors)
clfnn.fit(X_train, y_train)

y_scorenntr = clfnn.predict_proba(X_test)[: ,1]

ROC(y_test,y_scorenntr)
```

Area under the ROC curve : 0.606045



In [94]:

```
from sklearn import neighbors
n_neighbors = 15
clfnnad = neighbors.KNeighborsClassifier(n_neighbors)
clfnnad.fit(X, y)

y_scorenn = clfnnad.predict_proba(q)[: ,1]
```

In [95]:

```
# np.savetxt('K:/Mckinsey hackathon/nearest neighbor.txt',y_scorenn, delimiter="\t")
```